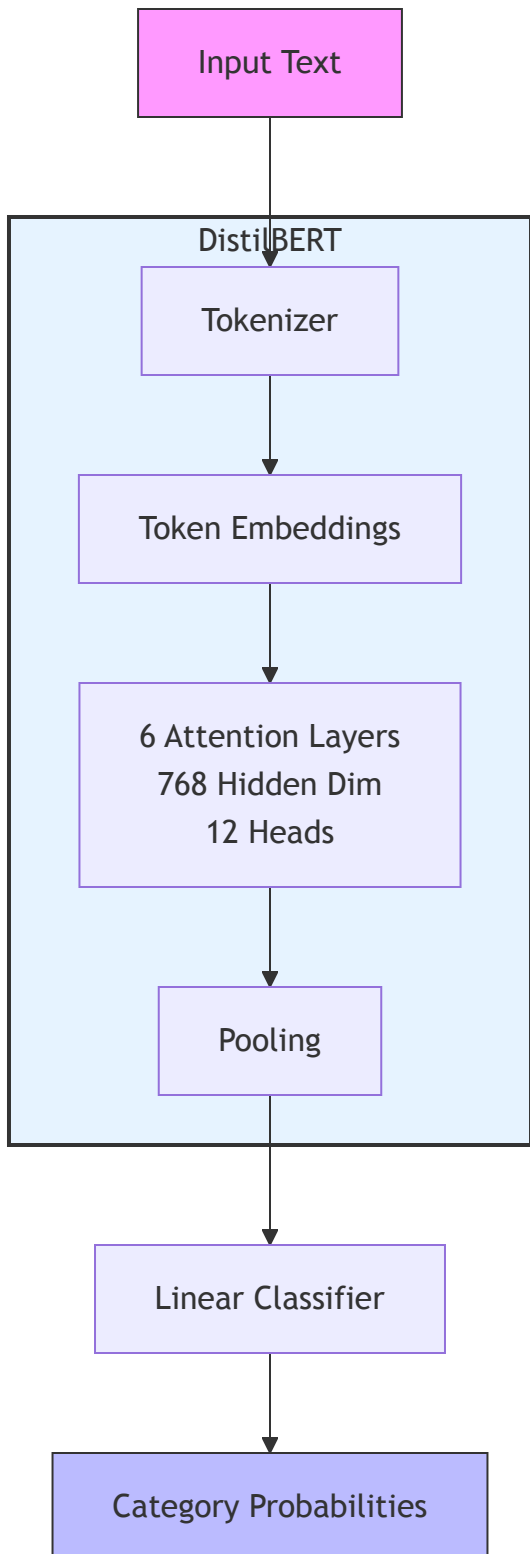


Child Protection Case Classification

Model

Model Architecture



1. Overview

Model Purpose

Automated classification of child protection case narratives into standardized categories for:

- Case prioritization
- Resource allocation
- Trend analysis

Version Information

| Component | Version |
|---------------|----------|
| Model Weights | 1.0.0 |
| Tokenizer | 1.0.0 |
| Training Data | v2024-06 |

2. Data Generation

Synthetic Data Pipeline

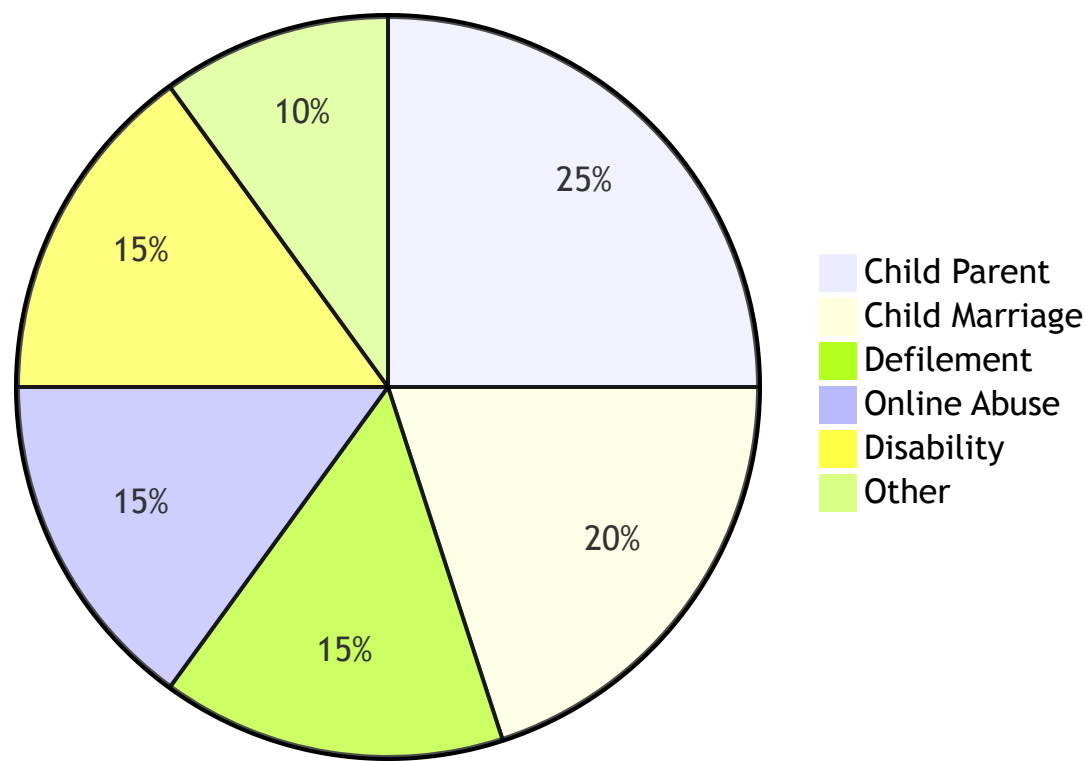
```
# Data generation pseudocode
1. Load Mistral-7B LLM
2. For each case:
  - Select random demographic profile
  - Generate narrative using category-specific prompt
  - Apply quality filters
3. Export to CSV
```

Dataset Statistics

| Metric | Value |
|-----------------------|------------|
| Total Cases | 1,000 |
| Avg. Narrative Length | 120 tokens |

| Metric | Value |
|------------------------|-------|
| Categories | 6 |
| Category Distribution: | |

Case Categories



3. Model Development

Training Configuration

```
model: distilbert-base-uncased
max_length: 256
batch_size: 16
epochs: 4
learning_rate: 2e-5
optimizer: AdamW
```

Training Progress

| Epoch | Loss | Grad Norm | LR |
|-------|------|-----------|---------|
| 1 | 2.70 | 1.80 | 1.89e-5 |
| 2 | 1.76 | 3.75 | 9.15e-6 |
| 4 | 1.30 | 3.60 | 6.10e-7 |

Key Metrics:

- Final Loss: 1.30
- Training Time: 4m28s
- Throughput: 9.6 samples/sec

4. Evaluation

Test Set Performance

| Metric | Score |
|-----------|-------|
| Accuracy | 0.88 |
| Precision | 0.86 |
| Recall | 0.85 |
| F1 Score | 0.85 |

Confusion Matrix:

| | | Predicted | | | |
|--------|---|-----------|----|----|------|
| | | A | B | C | D |
| Actual | A | [[85 | 5 | 3 | 7] |
| | B | 2 | 78 | 4 | 1] |
| | C | 1 | 4 | 82 | 3] |
| | D | 3 | 2 | 5 | 80]] |

5. Usage

Inference Example

```
from transformers import pipeline

classifier = pipeline(
    "text-classification",
    model="./case_category_model",
    tokenizer="distilbert-base-uncased"
)

case_text = "15yo female coerced into sending explicit images..."
result = classifier(case_text)
# Output: {'label': 'Online Abuse', 'score': 0.92}
```

API Specification

Endpoint: /classify-case

Input:

```
{
  "text": "case narrative...",
  "threshold": 0.7
}
```

Output:

```
{
  "prediction": "Child Marriage",
  "confidence": 0.91,
  "warning": "Requires urgent attention"
}
```

6. Limitations

Performance Boundaries

- Minimum text length: 50 characters
- Accuracy by category:
 - High (>90%): Defilement, Child Marriage
 - Medium (80-90%): Online Abuse
 - Lower (<80%): Complex multi-issue cases

Ethical Guidelines

- Always maintain human oversight
- Audit predictions for bias monthly
- Do not use for:
 - Legal determinations
 - Automatic service denial
 - Unverified reporting

7. Maintenance

Version History

| Version | Date | Changes |
|---------|------------|-----------------|
| 1.0.0 | 2024-06-01 | Initial release |

Retraining Protocol

Monthly:

- Generate new synthetic cases
- Fine-tune with 10% new data

Quarterly:

- Full retraining
- Bias assessment

8. Appendices

Requirements

```
transformers==4.30.2
torch==2.0.1
datasets==2.13.1
```

Training Log Sample

| Epoch | Loss | LR |
|-------|--------|---------|
| ----- | | |
| 1 | 2.7003 | 1.89e-5 |
| 2 | 1.7621 | 9.15e-6 |
| 4 | 1.3005 | 6.10e-7 |