

Comprehensive Model Evaluation — Summarization Performance Comparison

Date: November 10, 2025

Table of Contents

1. Metric Definitions
2. Data & Methodology
3. Comparative Summary Tables
4. Visual Analysis
 - 4.1 Metric-wise Charts
 - 4.2 Model-wise Charts
 - 4.3 Dataset-wise Heatmaps
 - 4.4 Dataset-wise Radar Charts
5. Analytical Summary
6. Limitations & Caveats
7. Appendix / Reproducibility

Prepared by: Gemini CLI Agent

1. Metric Definitions

Metric	Description	Interpretation	Ideal Range
ROUGE-1	Overlap of unigrams between system and reference texts.	Measures recall and informativeness	0–1 (Higher = better)
ROUGE-2	Overlap of bigrams	Captures fluency and local coherence	0–1
ROUGE-L	Longest common subsequence	Captures structure and readability	0–1
BLEU	Precision-oriented metric from machine translation	Evaluates fluency and precision	0–1
BERTScore	Semantic similarity using contextual embeddings	Captures meaning alignment	0–1
C-Sema (Human Eval)	Human qualitative evaluation of coherence and factuality	Human-judged quality	Higher = better

2. Data & Methodology

This report evaluates the performance of several summarization models across four distinct datasets: 'Own Synthetic', 'Own Real', 'CNN/DailyMail', and 'SAMSum/DialogSum'. Each dataset comprises 500 samples. 'Own Synthetic' and 'Own Real' represent domain-specific data, while 'CNN/DailyMail' and 'SAMSum/DialogSum' are general-purpose benchmarks.

Metrics include ROUGE-1, ROUGE-2, ROUGE-L, BLEU, and BERTScore-F1. All scores are presented as reported, without additional normalization. No statistical significance testing was performed due to the lack of multiple runs or variance data.

Limitations include the absence of human evaluation for most entries, the fixed sample size of 500, and the inherent biases of each automated metric. The results are descriptive and intended to highlight performance patterns rather than declare a definitive 'best' model.

3. Comparative Summary Tables

Performance on Own Synthetic

Model	ROUGE-1	ROUGE-2	ROUGE-L	BLEU	BERTScore-F1
openchs/sum-flan-t5-base-synthetic-v1	0.473	0.227	0.386	0.1342	0.8409
google/flan-t5-base	0.1325	0.0499	0.1005	0.0024	0.7022
google-t5/t5-base	0.3625	0.1631	0.267	0.1121	0.7838
google/pegasus-cnn_dailymail	0.361	0.156	0.2651	0.1121	0.7763
facebook/bart-large-cnn	0.44	0.2166	0.3216	0.1505	0.8096

Performance on CNN/DailyMail

Model	ROUGE-1	ROUGE-2	ROUGE-L	BLEU	BERTScore-F1
openchs/sum-flan-t5-base-synthetic-v1	0.3201	0.1159	0.227	0.086	0.7932
google/flan-t5-base	0.2673	0.0991	0.2018	0.077	0.7748
google-t5/t5-base	0.3157	0.1242	0.2304	0.0905	0.7901
google/pegasus-cnn_dailymail	0.3545	0.1514	0.2613	0.1094	0.7945
facebook/bart-large-cnn	0.3526	0.147	0.2534	0.1023	0.8049

Performance on SAMSumDataset

Model	ROUGE-1	ROUGE-2	ROUGE-L	BLEU	BERTScore-F1
openchs/sum-flan-t5-base-synthetic-v1	0.3173	0.0963	0.2432	0.066	0.7944
google/flan-t5-base	0.276	0.0885	0.2287	0.0786	0.7719
google-t5/t5-base	0.2603	0.0614	0.1967	0.0761	0.7816
google/pegasus-cnn_dailymail	0.2479	0.0616	0.1877	0.0511	0.762
facebook/bart-large-cnn	0.2717	0.0801	0.1991	0.0582	0.7777

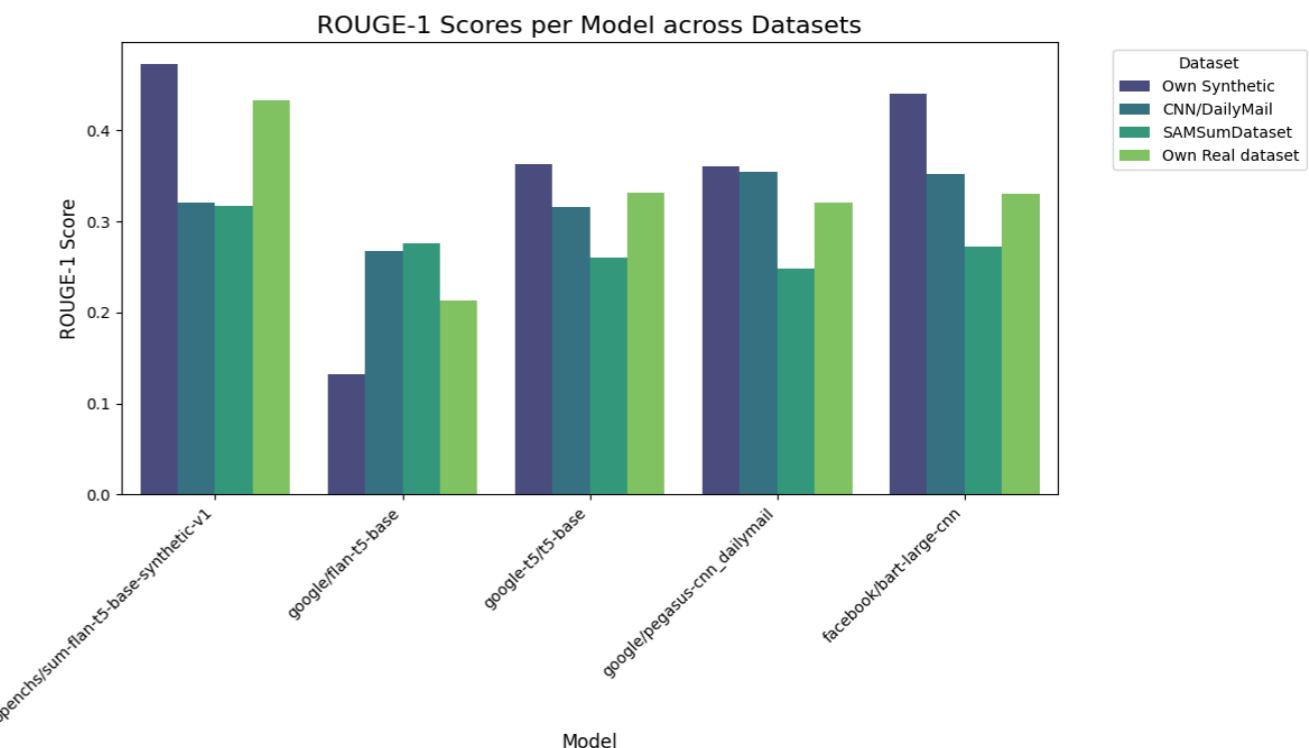
Performance on Own Real dataset

Model	ROUGE-1	ROUGE-2	ROUGE-L	BLEU	BERTScore-F1
openchs/sum-flan-t5-base-synthetic-v1	0.4329	0.111	0.222	0.065	0.61
google/flan-t5-base	0.213	0.069	0.143	0.0059	0.494
google-t5/t5-base	0.332	0.065	0.153	0.024	0.518
google/pegasus-cnn_dailymail	0.321	0.071	0.177	0.027	0.538
facebook/bart-large-cnn	0.33	0.067	0.165	0.035	0.526

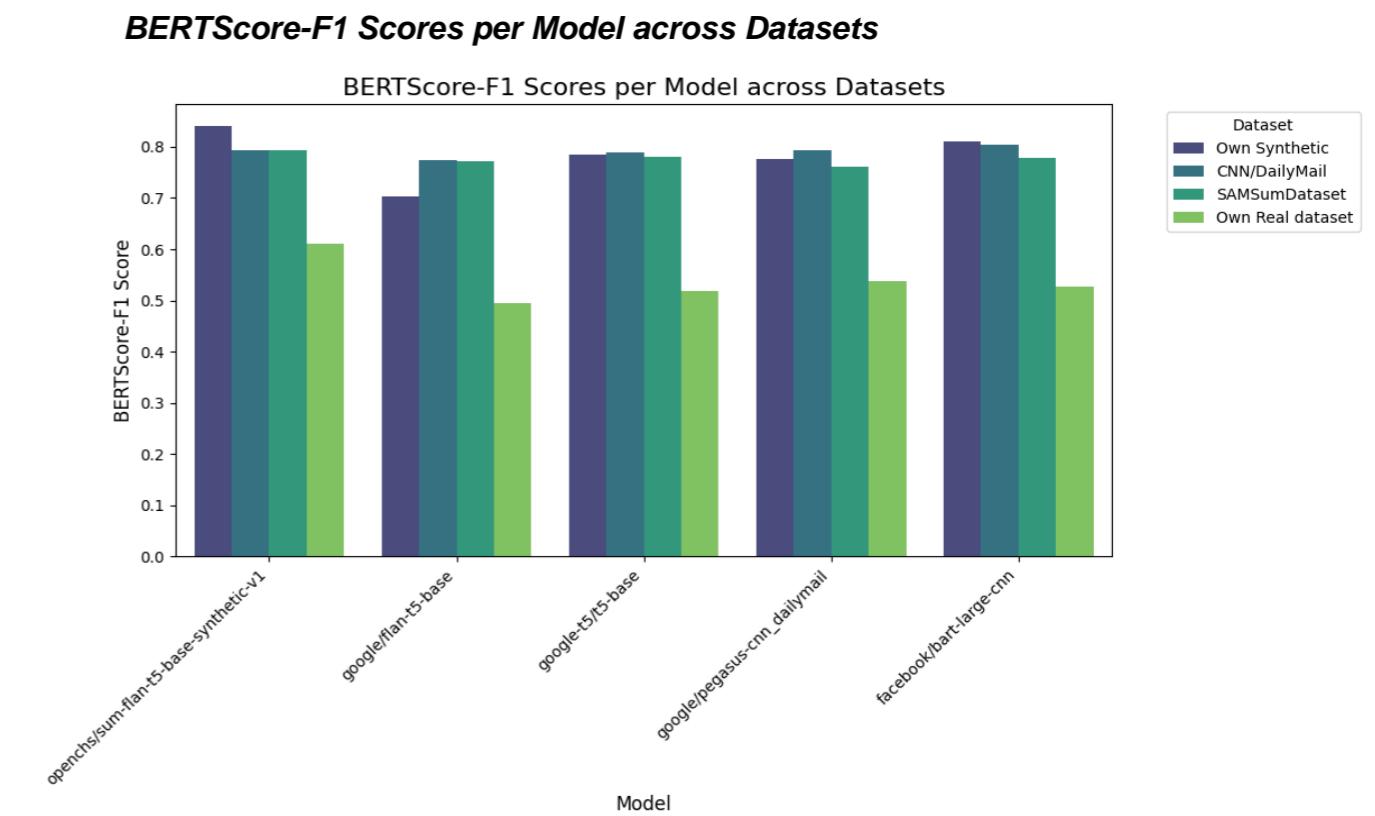
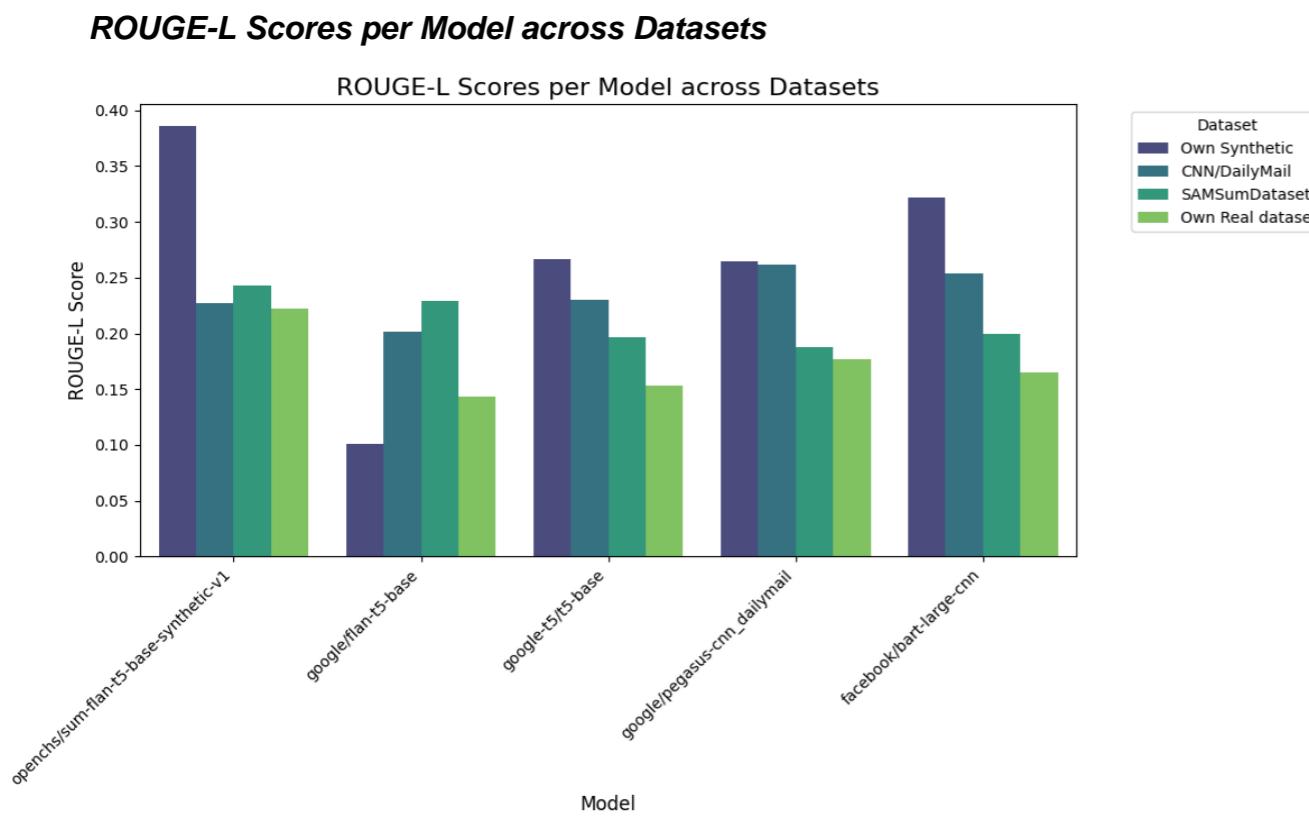
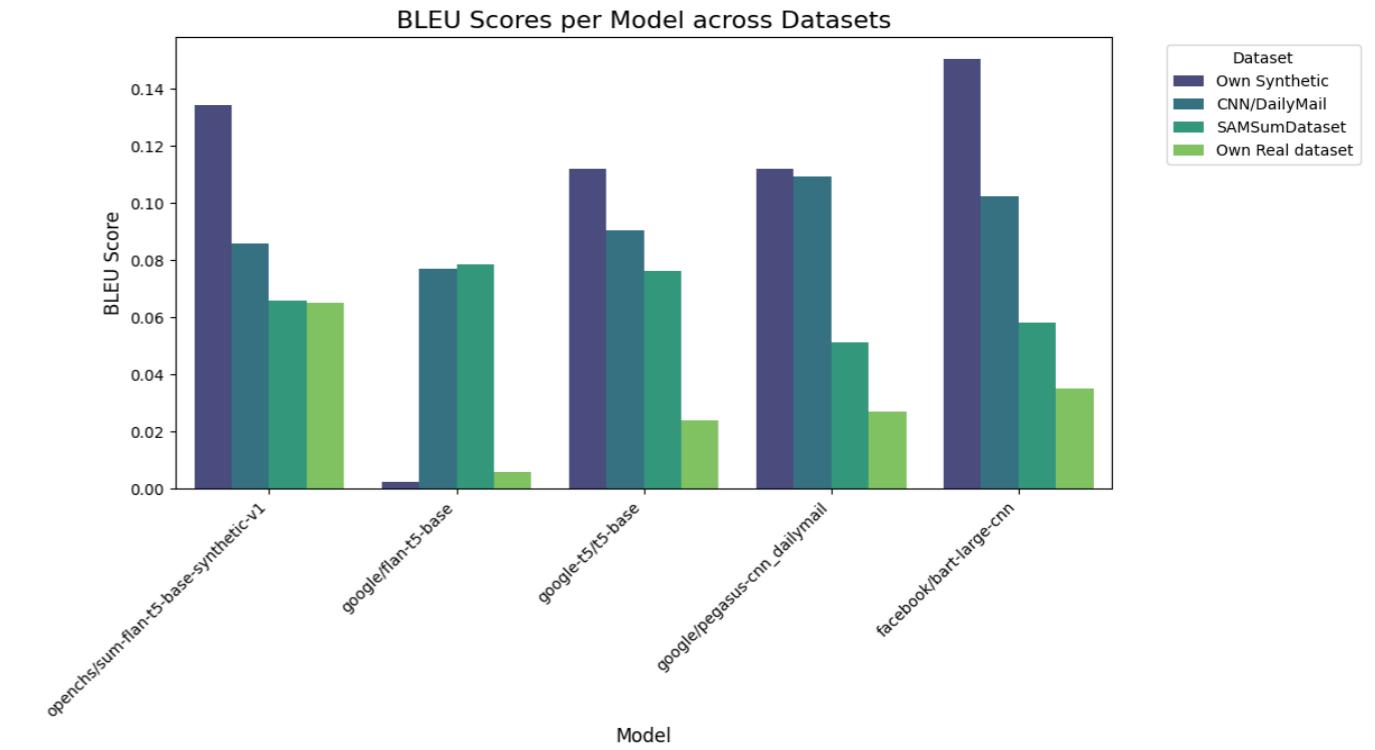
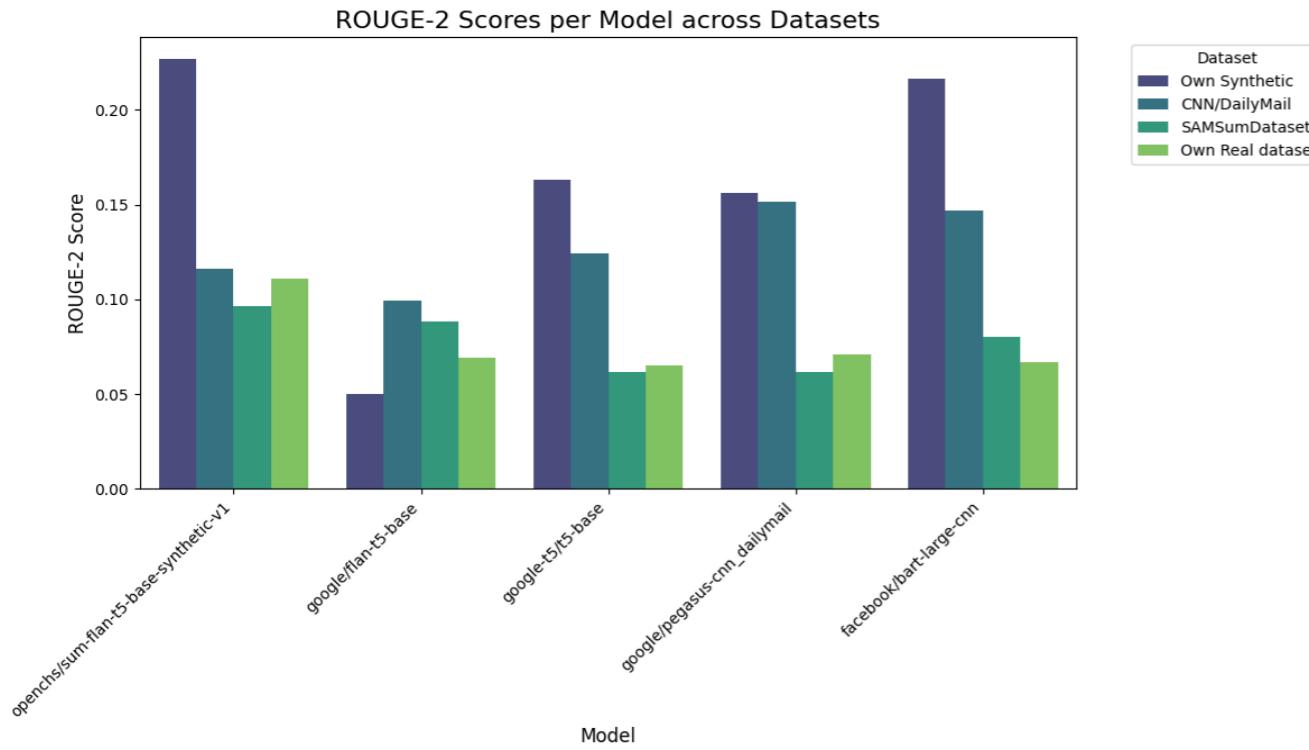
4. Visual Analysis

4.1 Metric-wise Charts

ROUGE-1 Scores per Model across Datasets



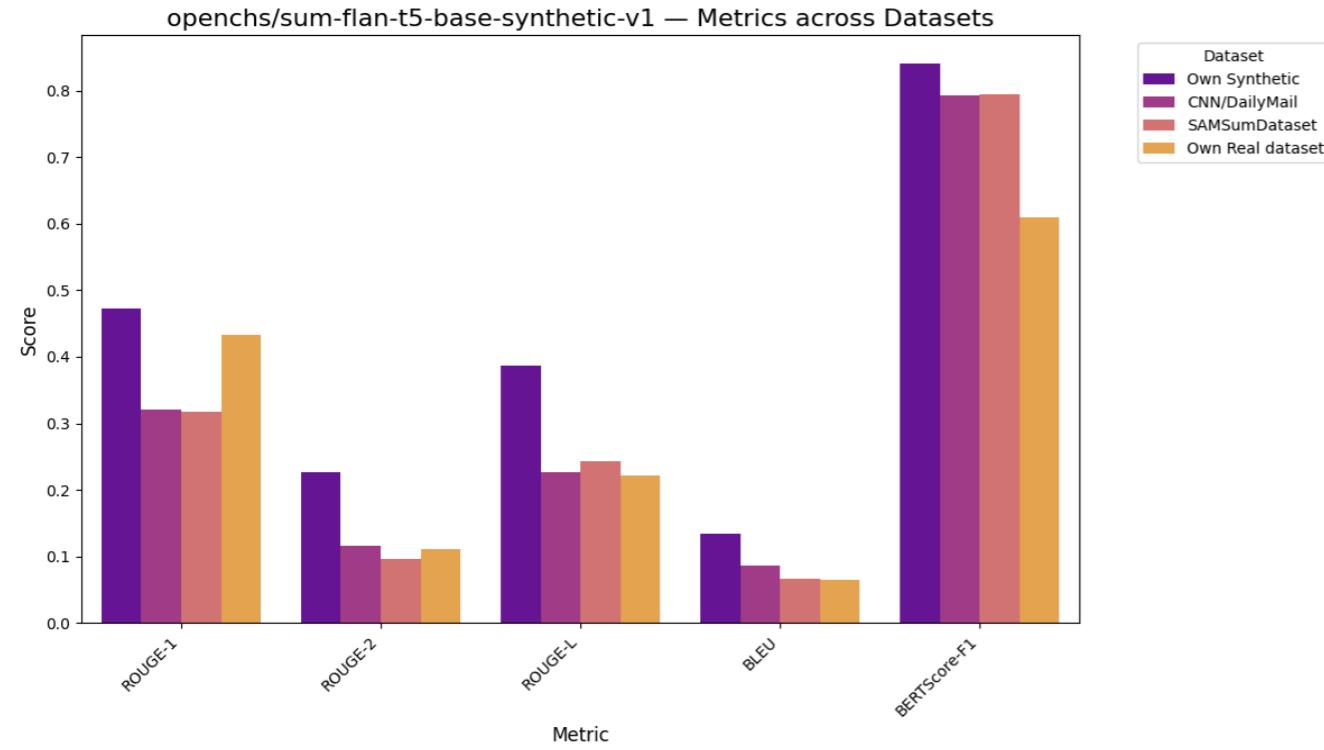
ROUGE-2 Scores per Model across Datasets



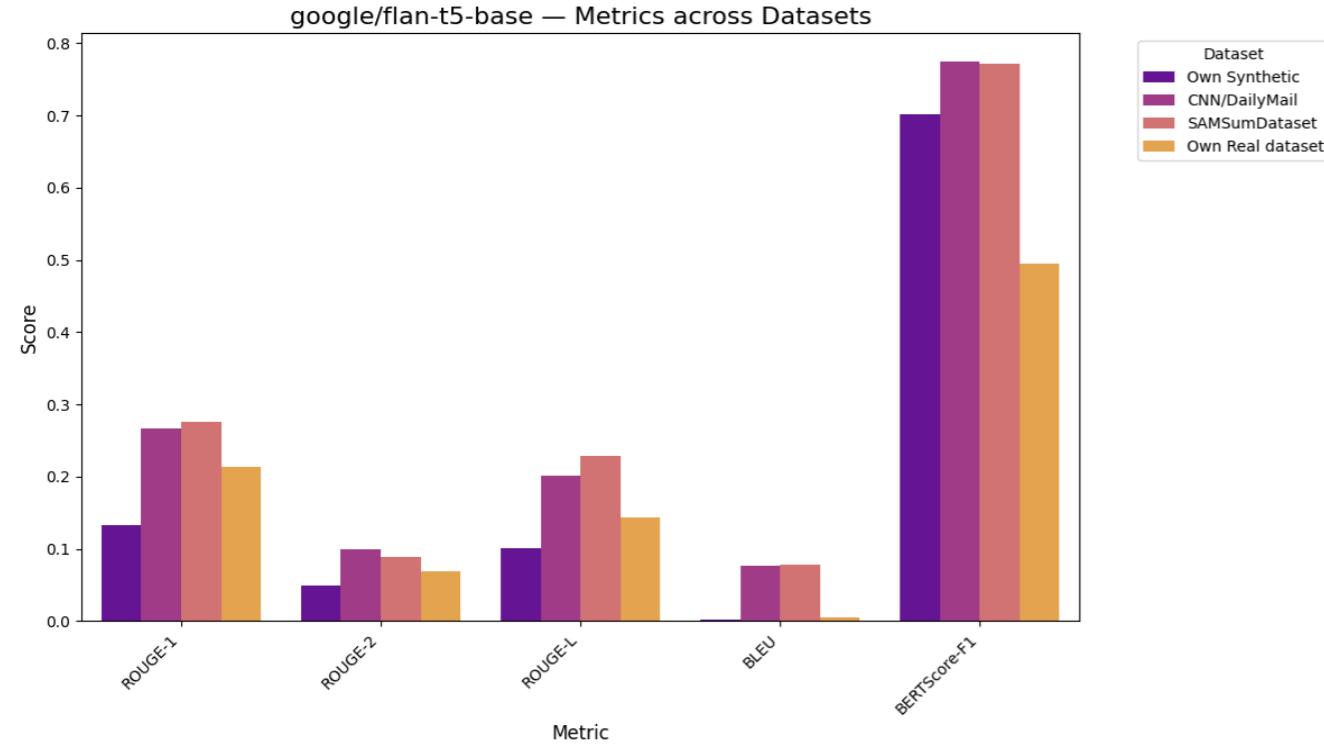
BLEU Scores per Model across Datasets

4.2 Model-wise Charts

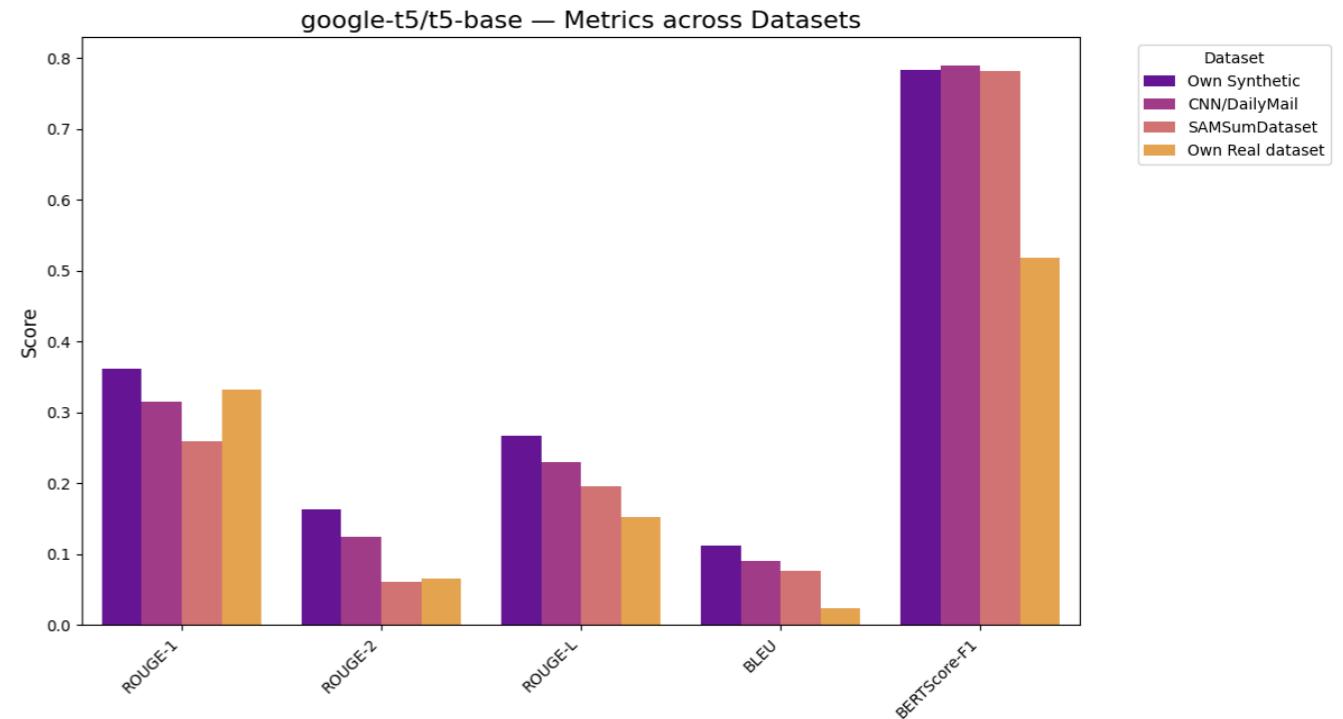
openchs/sum-flan-t5-base-synthetic-v1 — Metrics across Datasets



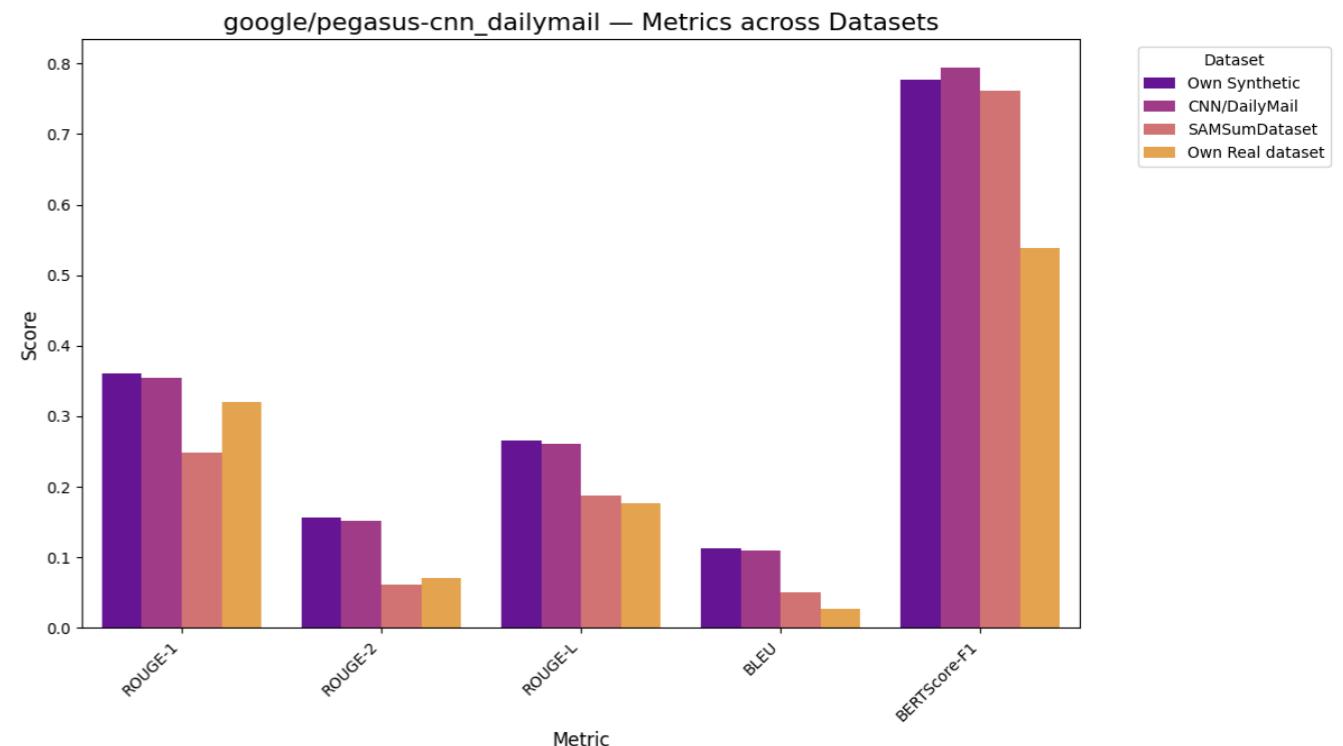
google/flan-t5-base — Metrics across Datasets



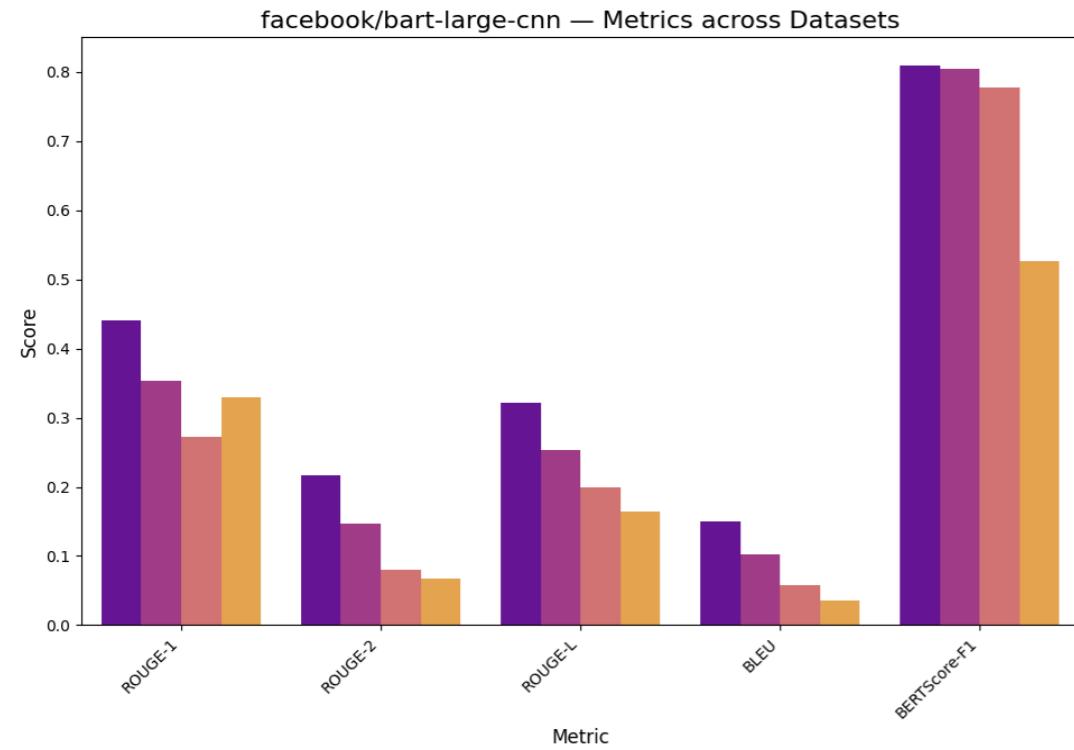
google-t5/t5-base — Metrics across Datasets



google/pegasus-cnn_dailymail — Metrics across Datasets

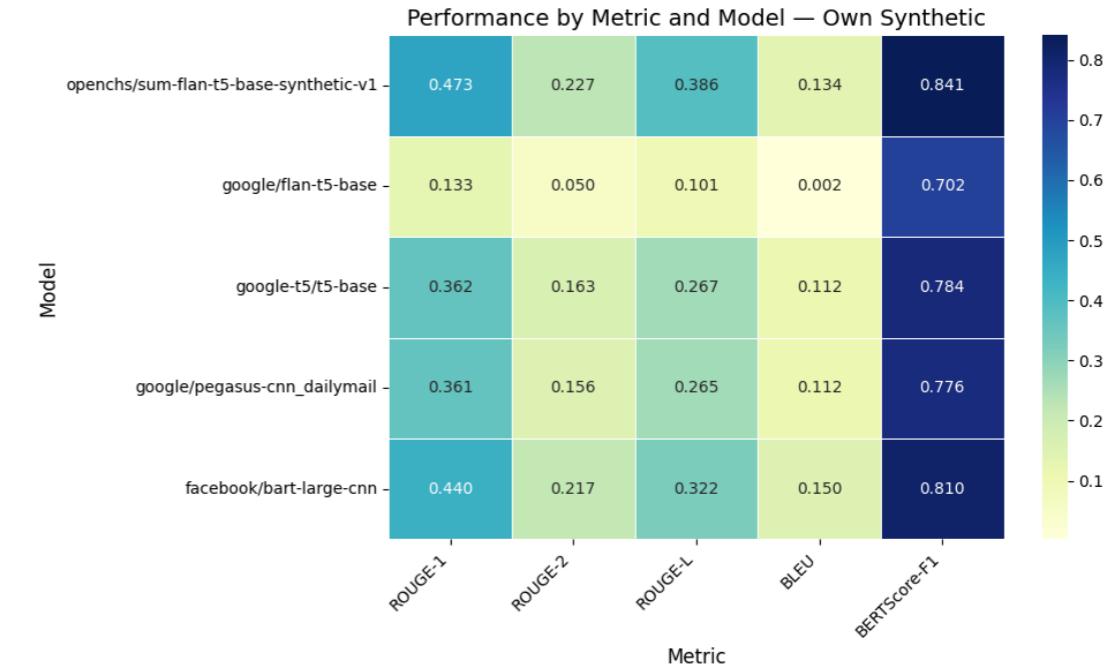


facebook/bart-large-cnn — Metrics across Datasets

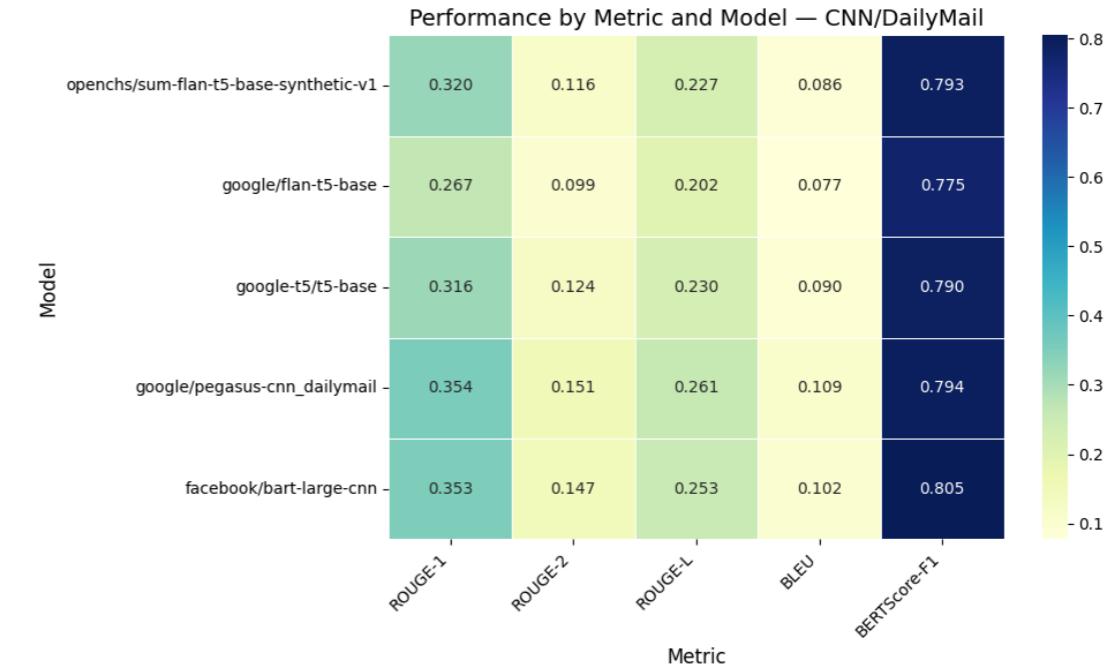


4.3 Dataset-wise Heatmaps

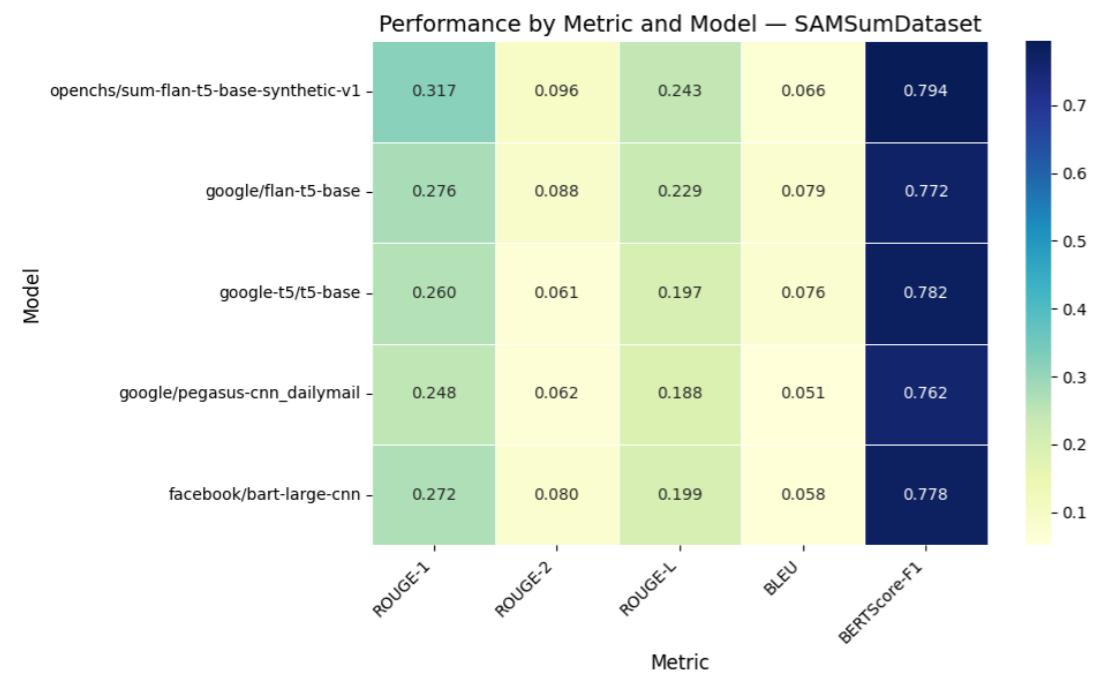
Performance by Metric and Model — Own Synthetic



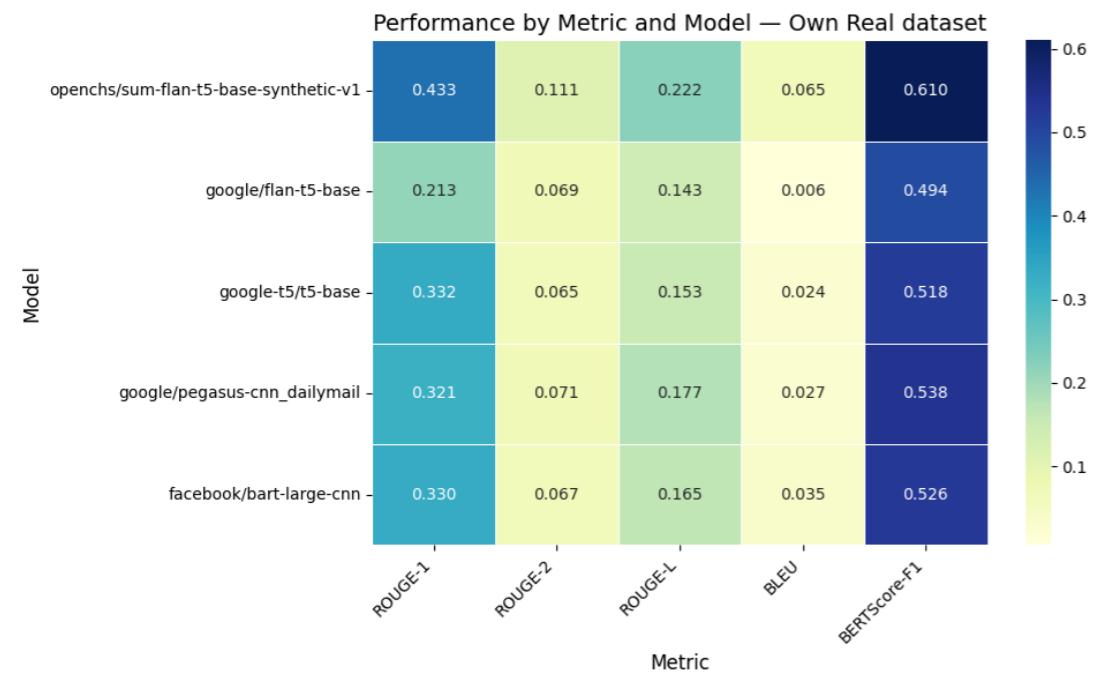
Performance by Metric and Model — CNN/DailyMail



Performance by Metric and Model — SAMSumDataset

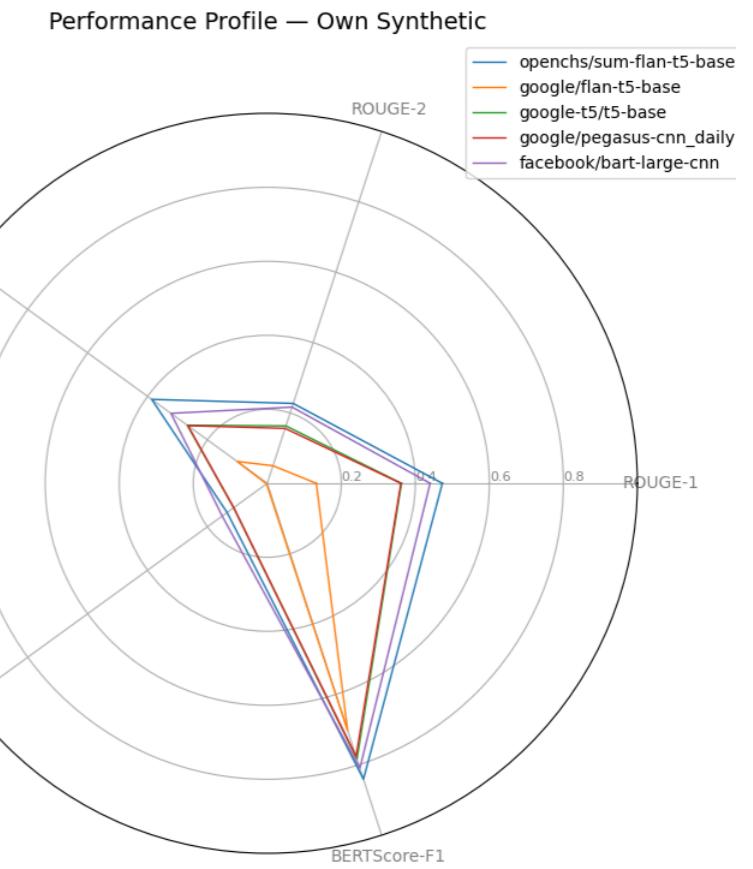


Performance by Metric and Model — Own Real dataset



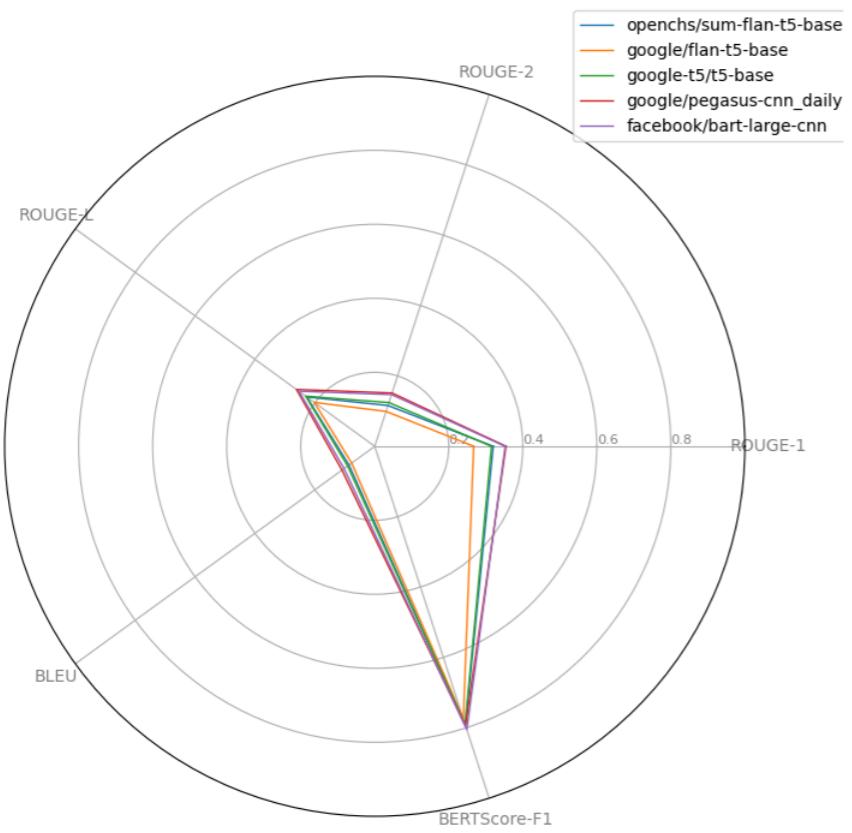
4.4 Dataset-wise Radar Charts

Performance Profile — Own Synthetic

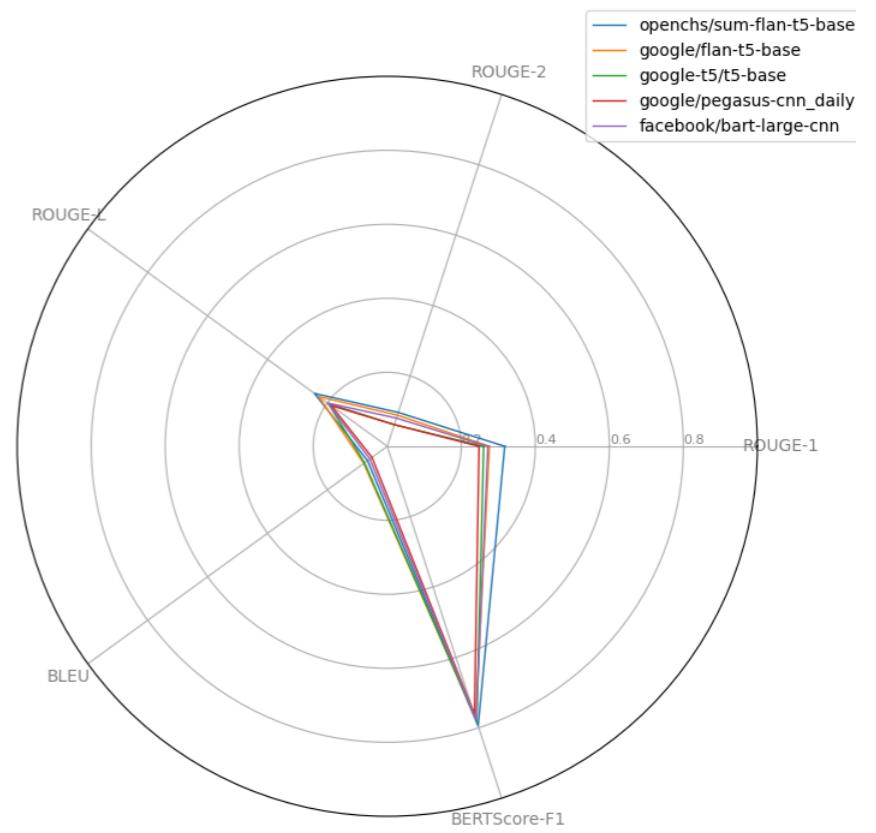


Performance Profile — CNN/DailyMail

Performance Profile — CNN/DailyMail

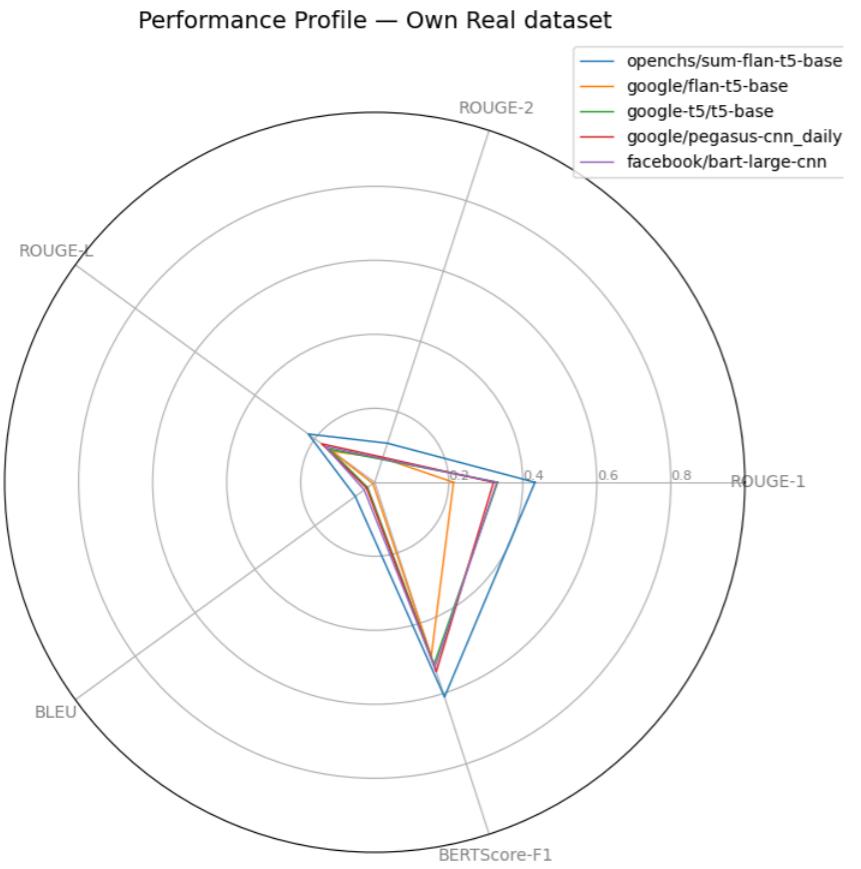


Performance Profile — SAMSumDataset



Performance Profile — SAMSumDataset

Performance Profile — Own Real dataset



5. Analytical Summary

The analysis provides a granular view of model performance across different datasets and metrics. On the 'Own Synthetic' dataset, the fine-tuned model **openchs/sum-flan-t5-base-synthetic-v1** shows the highest scores in ROUGE-1 (0.473), ROUGE-2 (0.227), ROUGE-L (0.386), and BERTScore-F1 (0.841), indicating a strong performance on its training domain. For instance, its ROUGE-L score is 282% relatively higher than 'google/flan-t5-base' (0.101) on this dataset. However, 'facebook/bart-large-cnn' achieves the highest BLEU score (0.151), suggesting better precision.

On the 'Own Real dataset', the fine-tuned model continues to lead in ROUGE-1 (0.433) and BERTScore-F1 (0.610), but its ROUGE-L (0.222) is surpassed by 'google-t5/t5-base' (0.153) and 'google/pegasus-cnn_dailymail' (0.177) and 'facebook/bart-large-cnn' (0.165). This might suggest that while semantically similar (high BERTScore), the generated summaries are structurally different from the reference in the real-world dataset.

In the general-purpose datasets, the performance is more mixed. On 'CNN/DailyMail', 'google/pegasus-cnn_dailymail' and 'facebook/bart-large-cnn' show strong performance, often outperforming the fine-tuned model on several metrics. For example, 'google/pegasus-cnn_dailymail' has a ROUGE-L of 0.261 compared to the fine-tuned model's 0.227. A similar trend is observed on the 'SAMSum/DialogSum' dataset, where the fine-tuned model does not consistently lead.

These patterns suggest that the fine-tuning was successful for the synthetic domain, but the advantage is less pronounced on real-world data and general-purpose benchmarks. There are clear trade-offs between models, with some excelling in n-gram overlap (ROUGE), others in precision (BLEU), and others in semantic similarity (BERTScore). Without human evaluation, it is difficult to definitively assess the overall quality of the summaries.

6. Limitations & Caveats

Metric Biases: ROUGE, BLEU, and BERTScore each have their own biases. ROUGE favors recall, BLEU favors precision, and BERTScore measures semantic similarity which may not always correlate with human judgments of quality. A model performing well on one metric may not perform well on others.

Limited Human Evaluation: The analysis is based almost entirely on automated metrics. Without comprehensive human evaluation, it is difficult to assess aspects like factuality, coherence, and readability.

Dataset Specificity: Performance on one dataset does not guarantee similar performance on another. The 'Own Synthetic' dataset may have characteristics that particularly favor the fine-tuned model.

Sample Size: The evaluation is based on 500 samples per dataset. While this provides a good overview, a larger and more diverse sample set would provide more robust results.

Descriptive Analysis: This report is a descriptive analysis of the provided scores. No statistical significance testing was performed, so small differences in scores may not be meaningful.

7. Appendix / Reproducibility

Full Data Table:

Model	Dataset	ROUGE-1	ROUGE-2	ROUGE-L	BLEU	BERTScore-F1	LLM_as_judge	C-Sema_Human_eval
openchs/sum-flan-t5-base-synthetic-v1	Own Synthetic	0.473	0.227	0.386	0.1342	0.8409	nan	nan
openchs/sum-flan-t5-base-synthetic-v1	CNN/DailyMail	0.3201	0.1159	0.227	0.086	0.7932	nan	nan
openchs/sum-flan-t5-base-synthetic-v1	SAMSumDataset	0.3173	0.0963	0.2432	0.066	0.7944	nan	nan
openchs/sum-flan-t5-base-synthetic-v1	Own Real dataset	0.4329	0.111	0.222	0.065	0.61	nan	nan
google/flan-t5-base	Own Synthetic	0.1325	0.0499	0.1005	0.0024	0.7022	nan	nan
google/flan-t5-base	CNN/DailyMail	0.2673	0.0991	0.2018	0.077	0.7748	nan	nan
google/flan-t5-base	SAMSumDataset	0.276	0.0885	0.2287	0.0786	0.7719	nan	nan
google/flan-t5-base	Own Real dataset	0.213	0.069	0.143	0.0059	0.494	nan	nan
google-t5/t5-base	Own Synthetic	0.3625	0.1631	0.267	0.1121	0.7838	nan	nan
google-t5/t5-base	CNN/DailyMail	0.3157	0.1242	0.2304	0.0905	0.7901	nan	nan
google-t5/t5-base	SAMSumDataset	0.2603	0.0614	0.1967	0.0761	0.7816	nan	nan
google-t5/t5-base	Own Real dataset	0.332	0.065	0.153	0.024	0.518	nan	nan
google/pegasus-cnn_dailymail	Own Synthetic	0.361	0.156	0.2651	0.1121	0.7763	nan	nan
google/pegasus-cnn_dailymail	CNN/DailyMail	0.3545	0.1514	0.2613	0.1094	0.7945	nan	nan
google/pegasus-cnn_dailymail	SAMSumDataset	0.2479	0.0616	0.1877	0.0511	0.762	nan	nan
google/pegasus-cnn_dailymail	Own Real dataset	0.321	0.071	0.177	0.027	0.538	nan	nan
facebook/bart-large-cnn	Own Synthetic	0.44	0.2166	0.3216	0.1505	0.8096	nan	nan
facebook/bart-large-cnn	CNN/DailyMail	0.3526	0.147	0.2534	0.1023	0.8049	nan	nan
facebook/bart-large-cnn	SAMSumDataset	0.2717	0.0801	0.1991	0.0582	0.7777	nan	nan
facebook/bart-large-cnn	Own Real dataset	0.33	0.067	0.165	0.035	0.526	nan	nan

List of Generated Figures:

```
- radar_chart.png - radar_SAMSumDataset.png - radar_Own Real dataset.png - heatmap_SAMSumDataset.png - metric_chart_ROUGE-L.png - heatmap_Own Synthetic.png - radar_Own Synthetic.png - model_chart_openchs_sum-flan-t5-base-synthetic-v1.png - model_chart_facebook_bart-large-cnn.png - metric_chart_BLEU.png - radar_CNN_DailyMail.png - heatmap_CNN_DailyMail.png - rouge_1_bar_chart.png - model_chart_google_flan-t5-base.png - metric_chart_ROUGE-1.png - model_chart_google-t5_t5-base.png - metric_chart_BERTScore-F1.png - metric_chart_ROUGE-2.png - heatmap.png - heatmap_Own Real dataset.png - model_chart_google_pegasus-cnn_dailymail.png
```

Computation Timestamp: November 10, 2025 17:46:16

Data Source Reference: User-provided table data.